



Domain Naming System

Lecture



Overview

- Introduction to the DNS
 - DNS Components
 - DNS Structure and Hierarchy
 - The DNS in Context
- 



DNS History (1)

- ARPANET utilized a central file HOSTS.TXT
 - Contains names to addresses mapping
 - Maintained by SRI's NIC (*Stanford-Research-Institute: Network-Information-Center*)
- Administrators email changes to NIC
 - NIC updates HOSTS.TXT periodically
- Administrators FTP (download) HOSTS.TXT



DNS History (2)

- As the system grew, HOSTS.TXT had problems with:
 - Scalability (traffic and load)
 - Name collisions
 - Consistency
- In **1984**, Paul Mockapetris released the first version (**RFCs 882** and **883**, superseded by **1034** and **1035** ...)



The DNS is...

- The “**Domain Name System**”
- What Internet users use **to reference anything by name** on the Internet.
- The mechanism by which Internet software translates **names to** attributes such as **addresses**

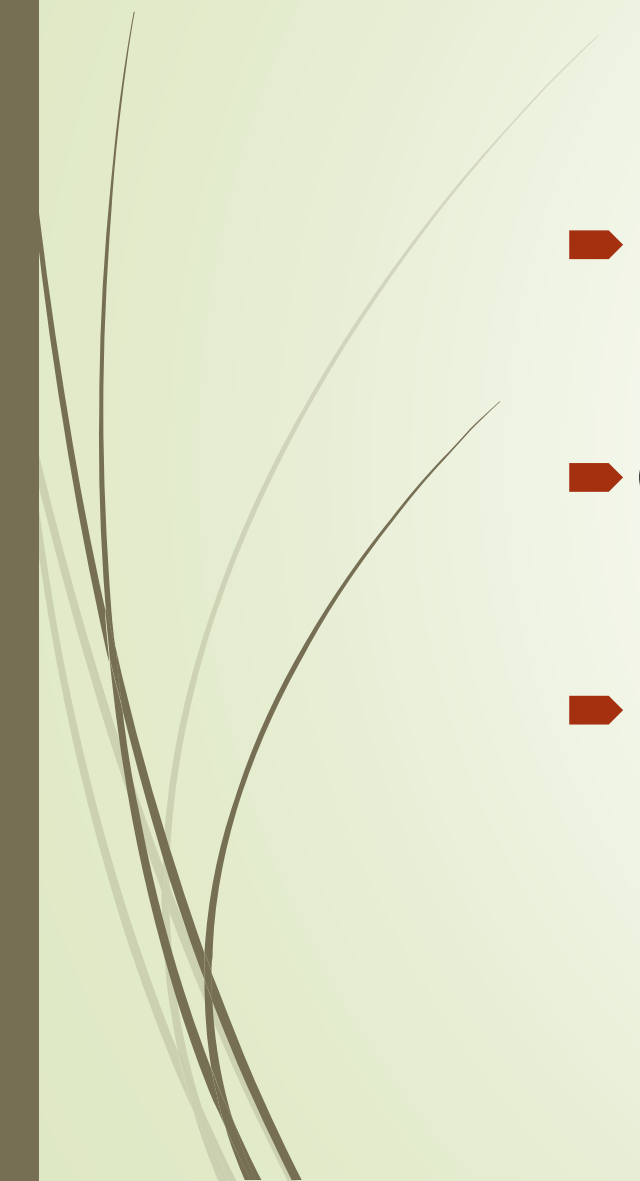


The DNS is also...

- A globally distributed, scalable, reliable **database**
- Comprised of three components
 - A “**name space**”
 - **Servers** making that name space available
 - **Resolvers** (clients) which query the servers about the name space



DNS as a Lookup Mechanism

- Users generally prefer names to numbers
 - Computers prefer numbers to names
 - DNS provides the mapping between the two
 - I have “x”, give me “y”
- 



DNS as a Database

- Keys to the database are “**domain names**”
 - www.test.com, 18.in-addr.arpa, 6.4.e164.arpa
- Over 200,000,000 domain names stored
- Each **domain** name contains **one or more attributes**
 - Known as “**resource records**”
- Each attribute **individually retrievable**



Global Distribution

- ➡ Data is **maintained locally**, but **retrievable globally**
 - ➡ No single computer has all DNS data
- ➡ **DNS lookups** can be performed by **any device**
- ➡ Remote DNS data is **locally cachable** to improve performance



Scalability

- No limit to the **size of the database**
- No limit to the **number of queries**
 - Tens of thousands of queries handled easily every second
- **Queries distributed** among masters, slaves, and caches



Reliability

- Data is replicated
 - Data from master is copied to multiple slaves(replicas)
- Clients can query
 - Master server
 - Any of the copies at slave servers
- Clients will typically query local caches
- DNS protocols can use either UDP or TCP
 - If UDP, DNS protocol handles retransmission, sequencing, etc.

Loose Coherency

- Each version of a subset of the database (a **zone**) has a **serial number**
 - The serial number is **incremented** on each **database change**
- **Changes** to the **master** copy of the database are **propagated to replicas** according to timing set by the zone administrator
- **Cached** data **expires** according to **timeout** set by **zone administrator**

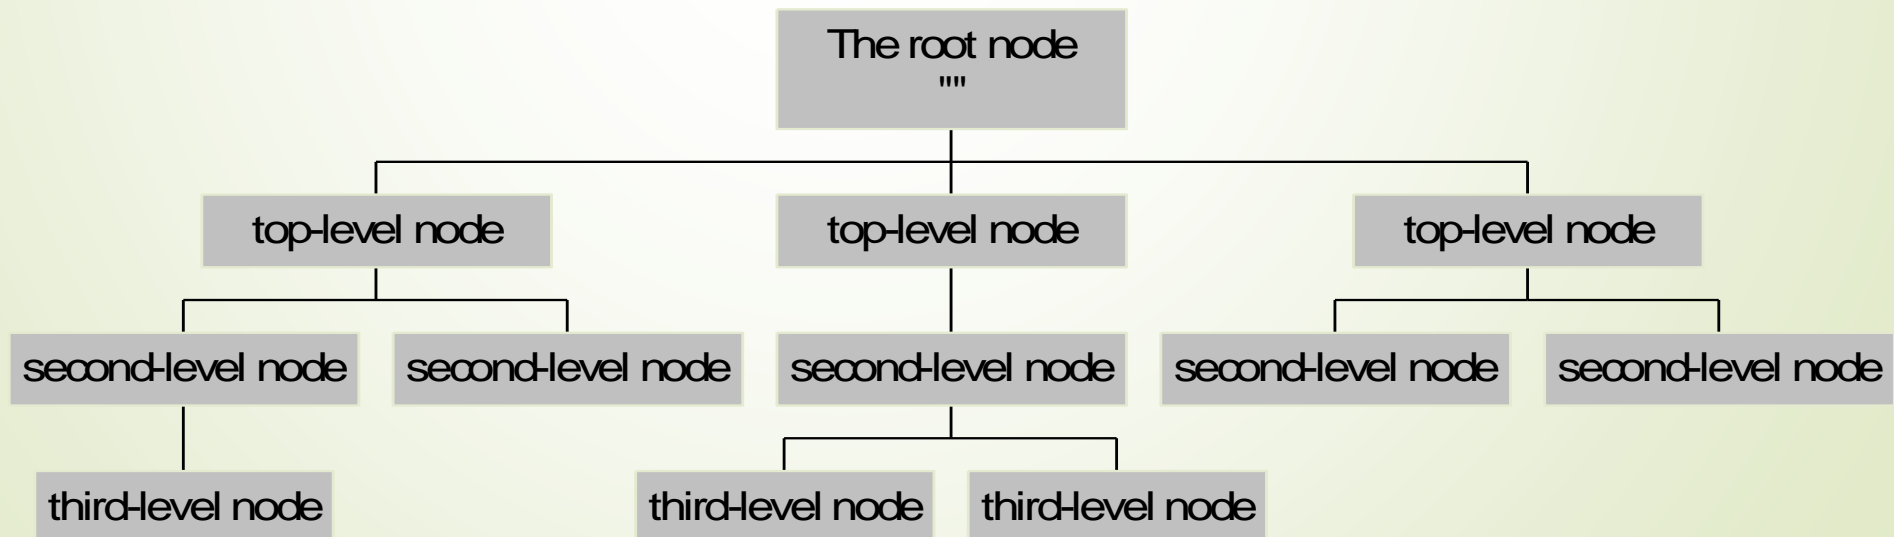


Dynamicity

- Database can be **updated dynamically**
 - Add/delete/modify of any record
 - **Only master** can be dynamically updated
- **Modification** of the master database **triggers replication**

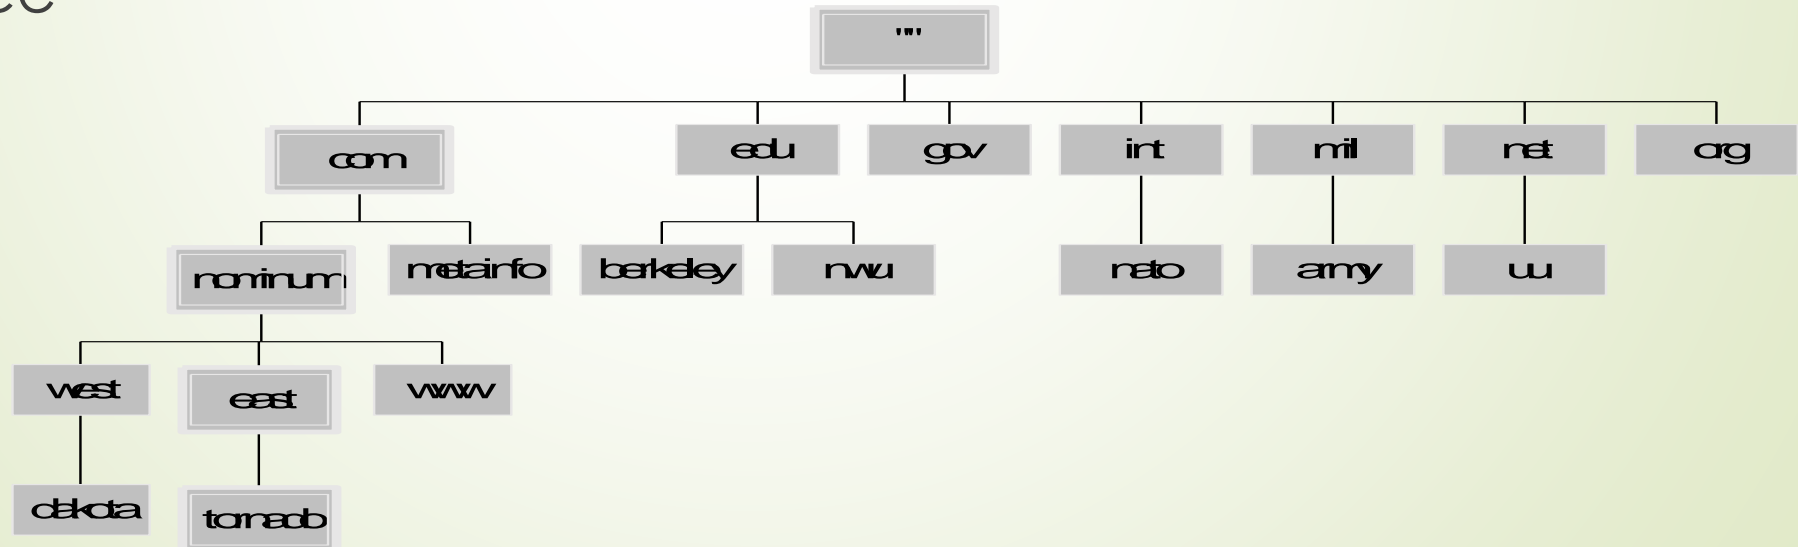
The Name Space

- The ***name space*** is the structure of the DNS database
 - An **inverted tree** with the **root node** at the top
- Each node has a label
 - The root node has a null label, written as ""



Domain Names

- ▶ A *domain name* is the sequence of labels from a node to the root, separated by dots (".")s, read left to right
 - ▶ The name space has a maximum depth of 127 levels
 - ▶ Domain names are limited to 255 characters in length
- ▶ A node's domain name identifies its position in the name space






Subdomains

- One domain is a subdomain of another if its domain name ends in the other's domain name
 - So *sales.nominum.com* is a subdomain of
 - *nominum.com* & *com*
 - *nominum.com* is a subdomain of *com*



Delegation

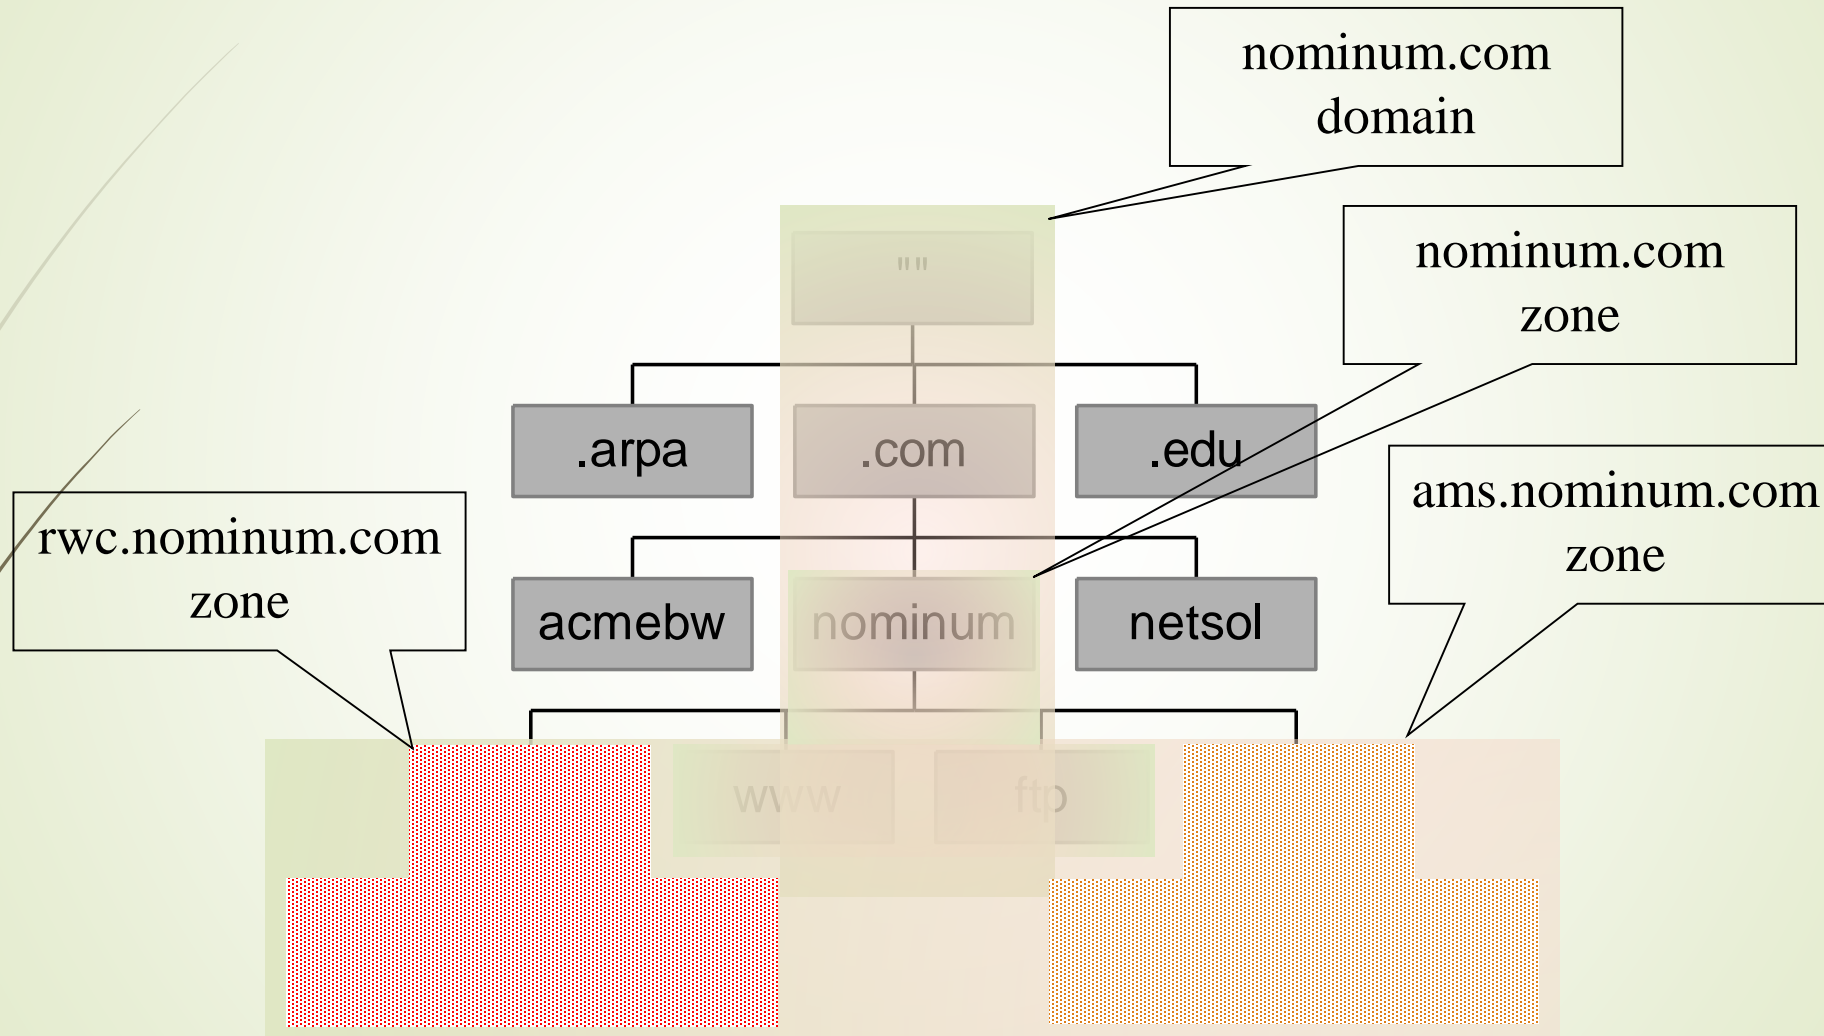
- Administrators can create subdomains to group hosts
 - According to geography, organizational affiliation etc.
 - An administrator of a domain can delegate responsibility for managing a subdomain to someone else
 - The parent domain retains links to the delegated subdomains
- 



Delegation Creates Zones

- Each time an administrator delegates a subdomain, a new unit of administration is created
 - The subdomain and its parent domain can now be administered independently
 - These units are called *zones*
 - The boundary between zones is a point of delegation in the name space
- Delegation is good: it is the key to scalability

Dividing a Domain into Zones






Name Servers

- ▶ Name servers store information about the name space in units called “zones”
 - ▶ The name servers that load a complete zone are said to “have authority for” or “be authoritative for” the zone
- ▶ Usually, more than one name server are **authoritative** for the same zone
 - ▶ This ensures redundancy and spreads the load
- ▶ Also, a single name server may be authoritative for many zones

Name Servers and Zones





Types of Name Servers

- Two main types of servers
 - Authoritative – maintains the data
 - Master – where the data is edited
 - Slave – where data is replicated to
 - Caching – stores data obtained from an authoritative server
- No special hardware necessary



Name Server Architecture

- You can think of a name server as part of:
 - *database server*, answering queries about the parts of the name space it knows about (i.e., is authoritative for),
 - *cache*, temporarily storing data it learns from other name servers, and
 - *agent*, helping resolvers and other name servers find data


Name Server Architecture

Name Server Process

Authoritative Data
(primary master and
slave zones)


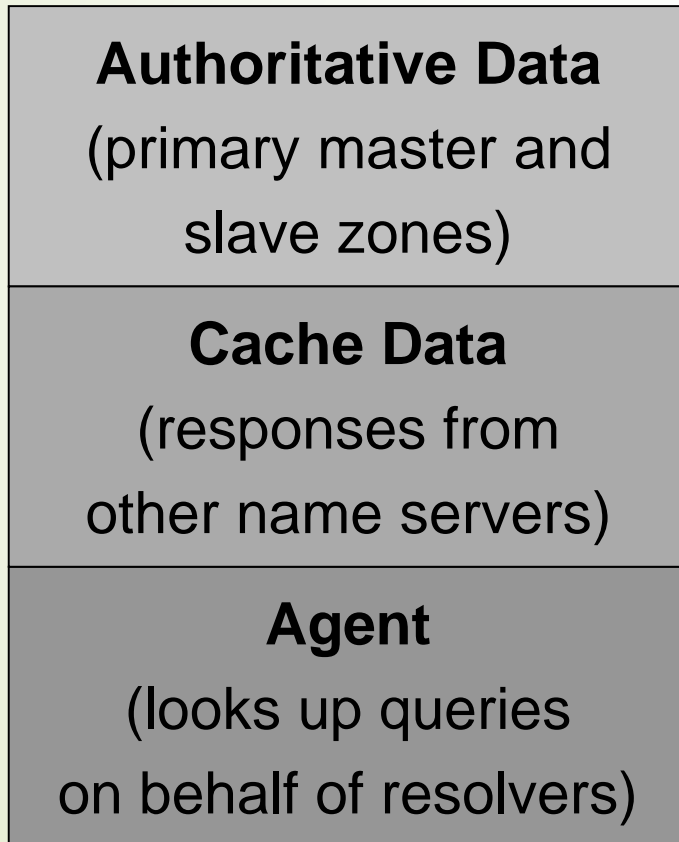
Cache Data
(responses from
other name servers)

Agent
(looks up queries
on behalf of resolvers)



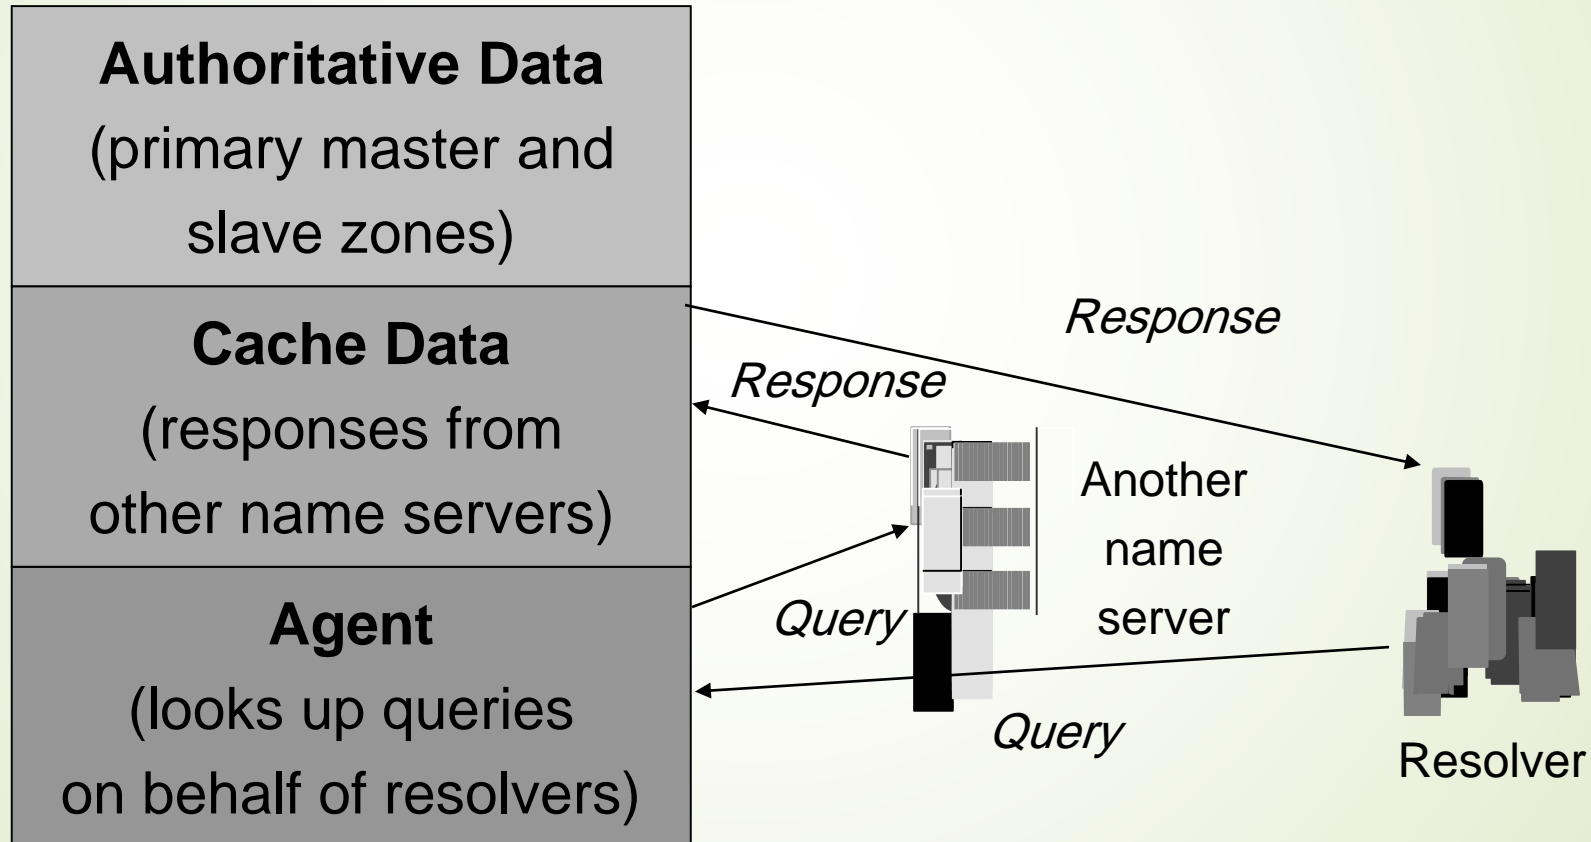
Authoritative Data

Name Server Process



Using Other Name Servers

Name Server Process



Cached Data

Name Server Process


Authoritative Data
(primary master and
slave zones)

Cache Data
(responses from
other name servers)

Agent
(looks up queries
on behalf of resolvers)

Response

Query



Resolver



Name Resolution

- *Name resolution* is the process by which resolvers and name servers cooperate to find data in the name space
- Closure mechanism for DNS?
 - Starting point: the names and IP addresses of the name servers for the root zone (the “root name servers”)
 - The root name servers know about the top-level zones and can tell name servers whom to contact for all TLDs

Name Resolution

- A DNS query has three parameters:

1. A domain name (e.g., *www.nominum.com*),
 - Remember, every node has a domain name!
2. A class (e.g., *IN*), and
3. A type (e.g., *A*)

<http://network-tools.com/nslook/>

- Upon receiving a query from a resolver, a name server

1. looks for the answer in its authoritative data and its cache
2. If step 1 fails, the answer must be looked up

The Resolution Process

- ▶ Let's look at the resolution process step-by-step:




`annie.west.sprockets.com`

`ping www.nominum.com.`

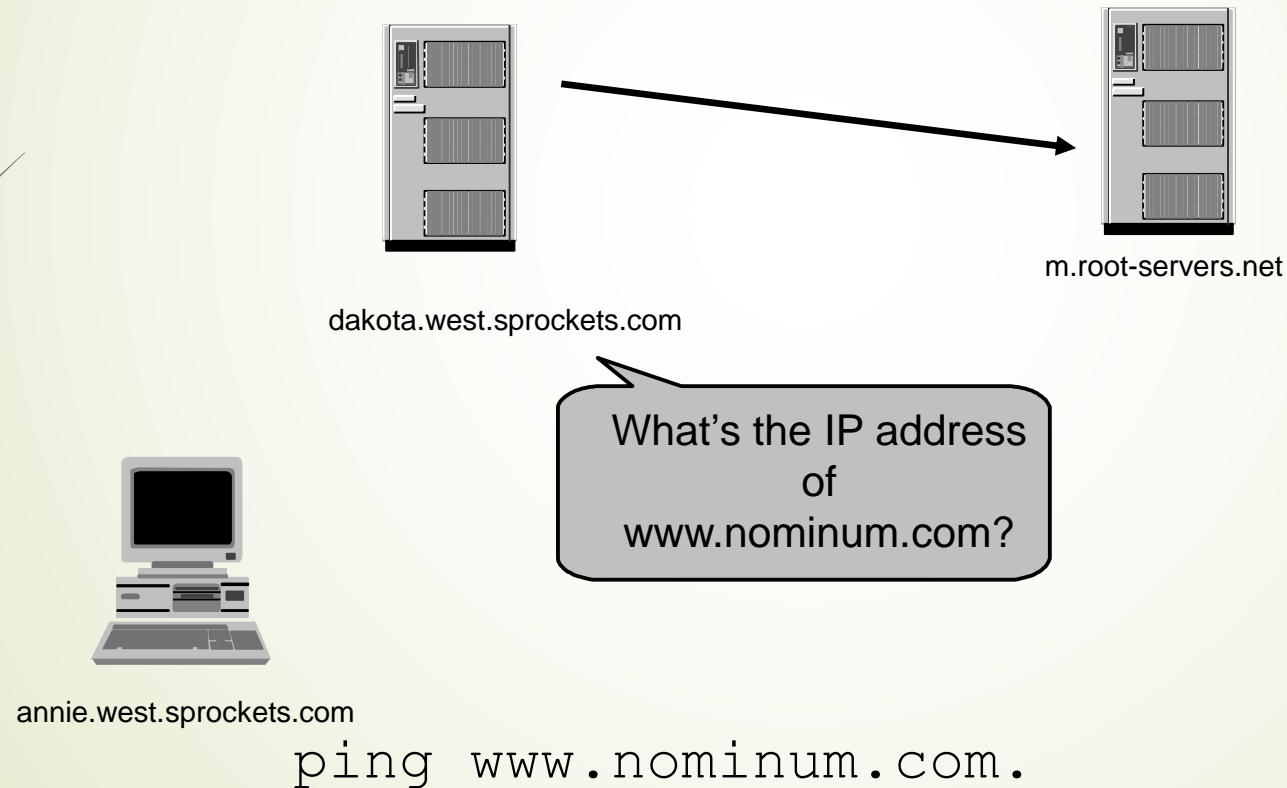
The Resolution Process

- The workstation *annie* asks its configured name server, *dakota*, for *www.nominum.com*'s address



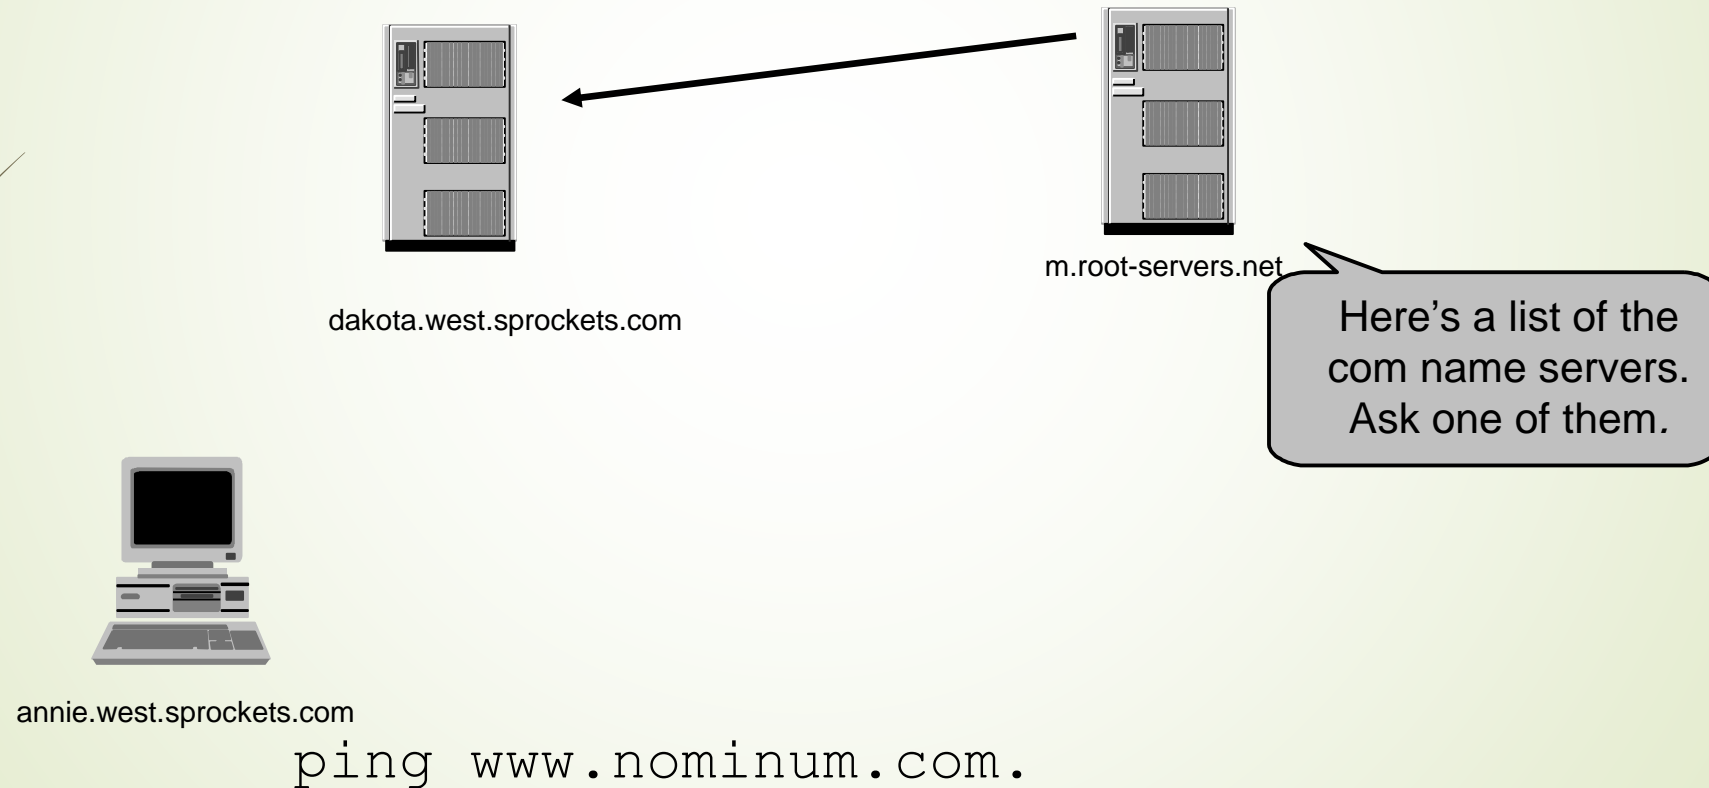
The Resolution Process

- The name server *dakota* asks a root name server, *m*, for *www.nominum.com*'s address



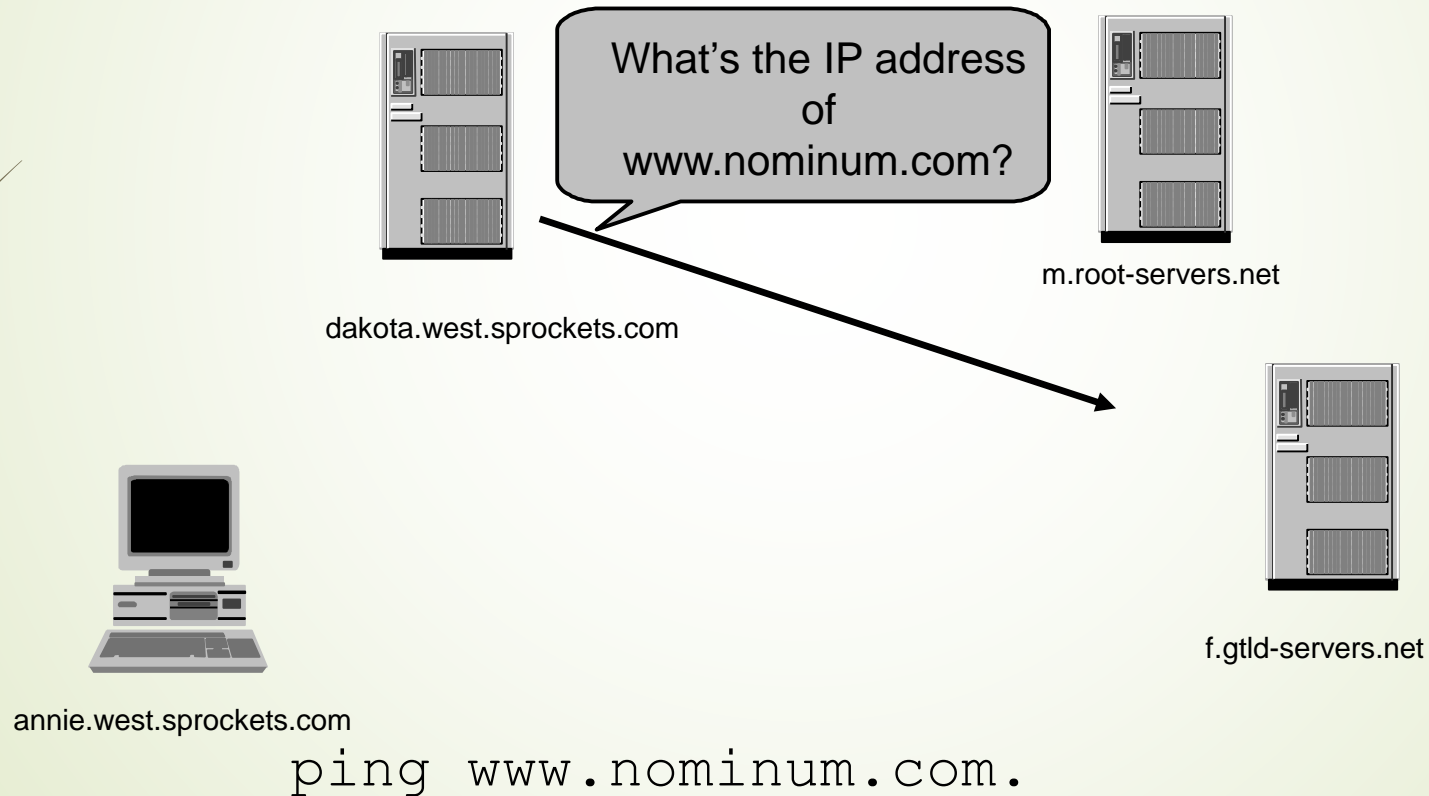
The Resolution Process

- The root server *m* refers *dakota* to the *com* name servers
- This type of response is called a “referral”




The Resolution Process

- ▶ The name server *dakota* asks a *com* name server, *f*, for *www.nominum.com*'s address



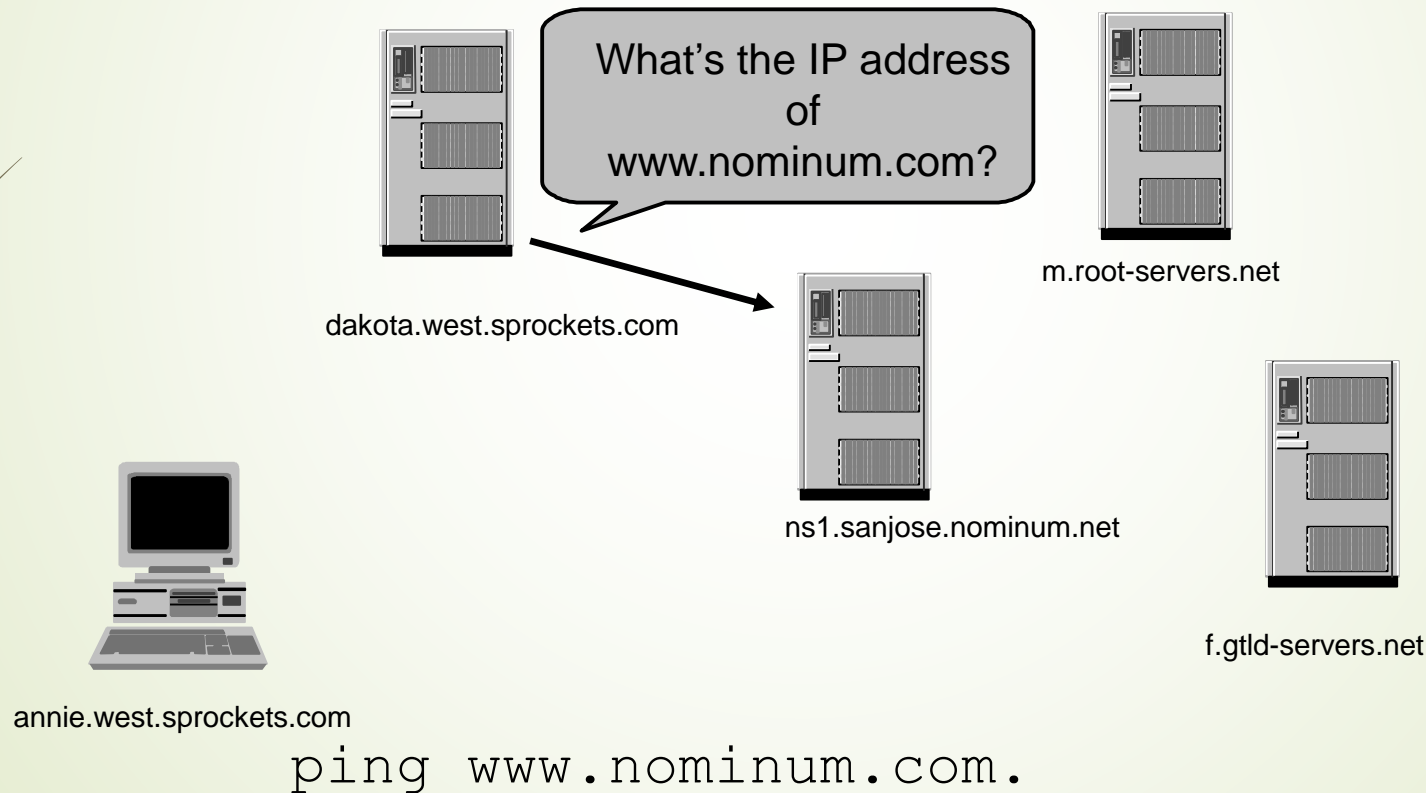
The Resolution Process

- The com name server *f* refers *dakota* to the *nominum.com* name servers



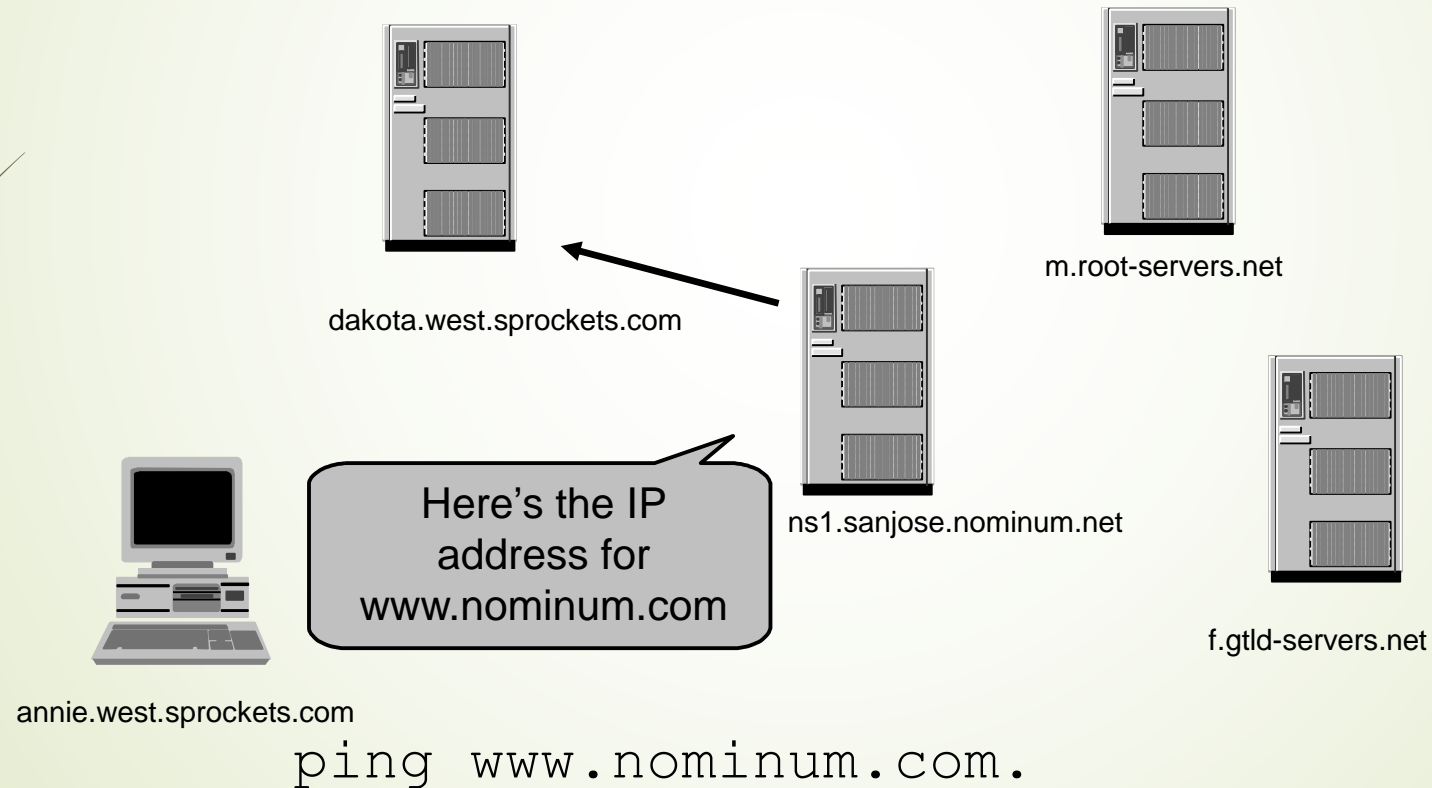
The Resolution Process

- The name server *dakota* asks a *nominum.com* name server, *ns1.sanjose*, for *www.nominum.com*'s address




The Resolution Process

- ▶ The *nominum.com* name server *ns1.sanjose* responds with *www.nominum.com*'s address



The Resolution Process

- The name server *dakota* responds to *annie* with *www.nominum.com*'s address



Resolution Process (Caching)

- After the previous query, the name server *dakota* now knows:
 - The names and IP addresses of the *com* name servers
 - The names and IP addresses of the *nominum.com* name servers
 - The IP address of *www.nominum.com*
- Let's look at the resolution process again

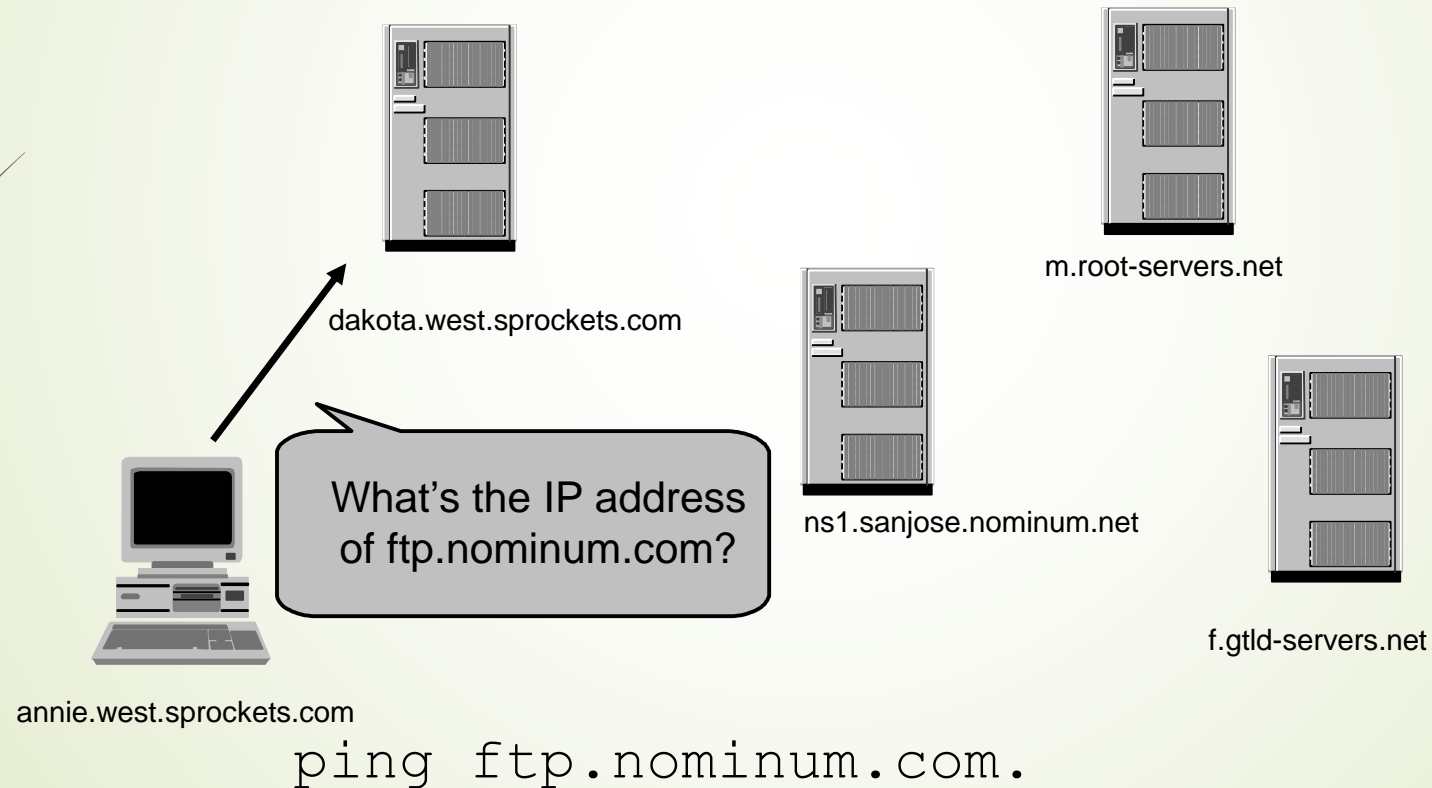


annie.west.sprockets.com

ping **ftp**.nominum.com.

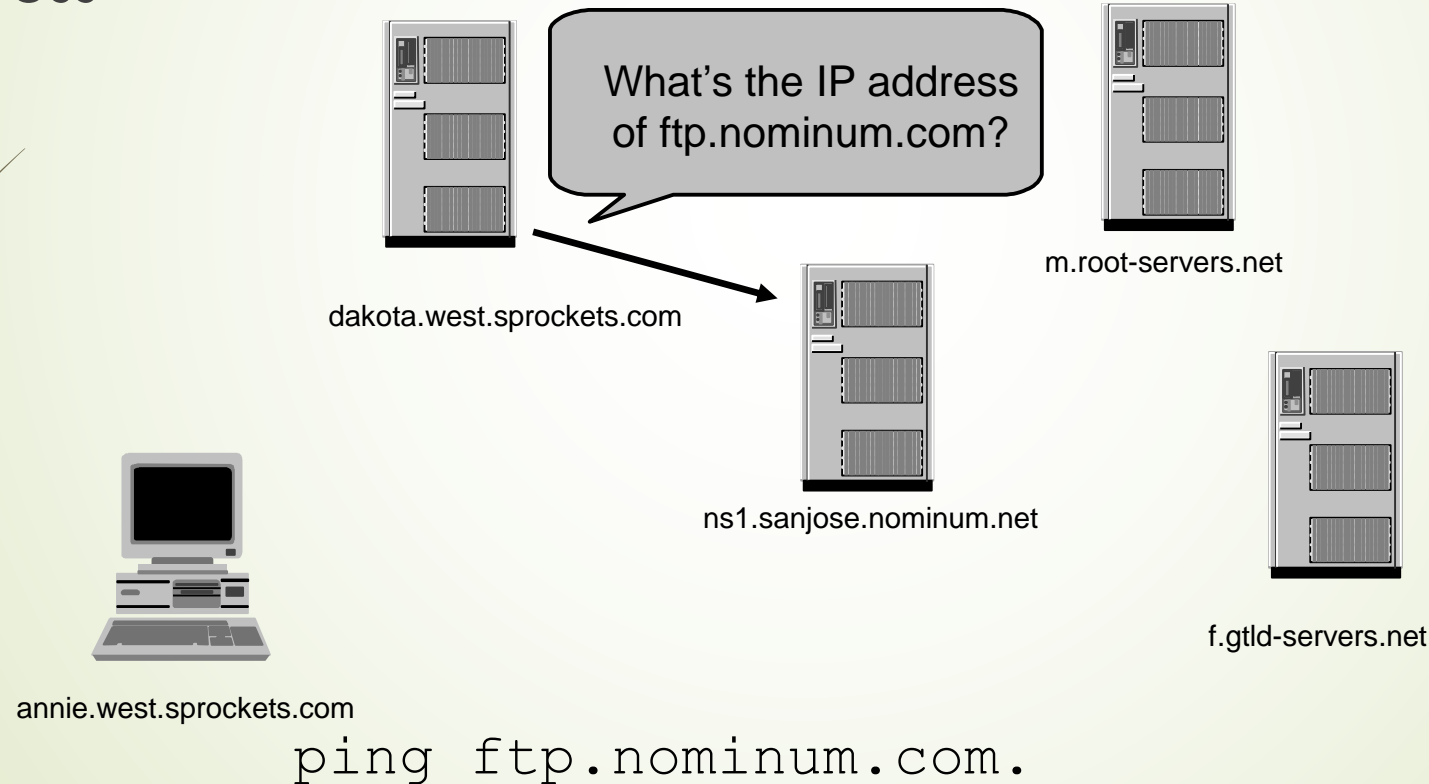
Resolution Process (Caching)

- The workstation *annie* asks its configured name server, *dakota*, for *ftp.nominum.com*'s address



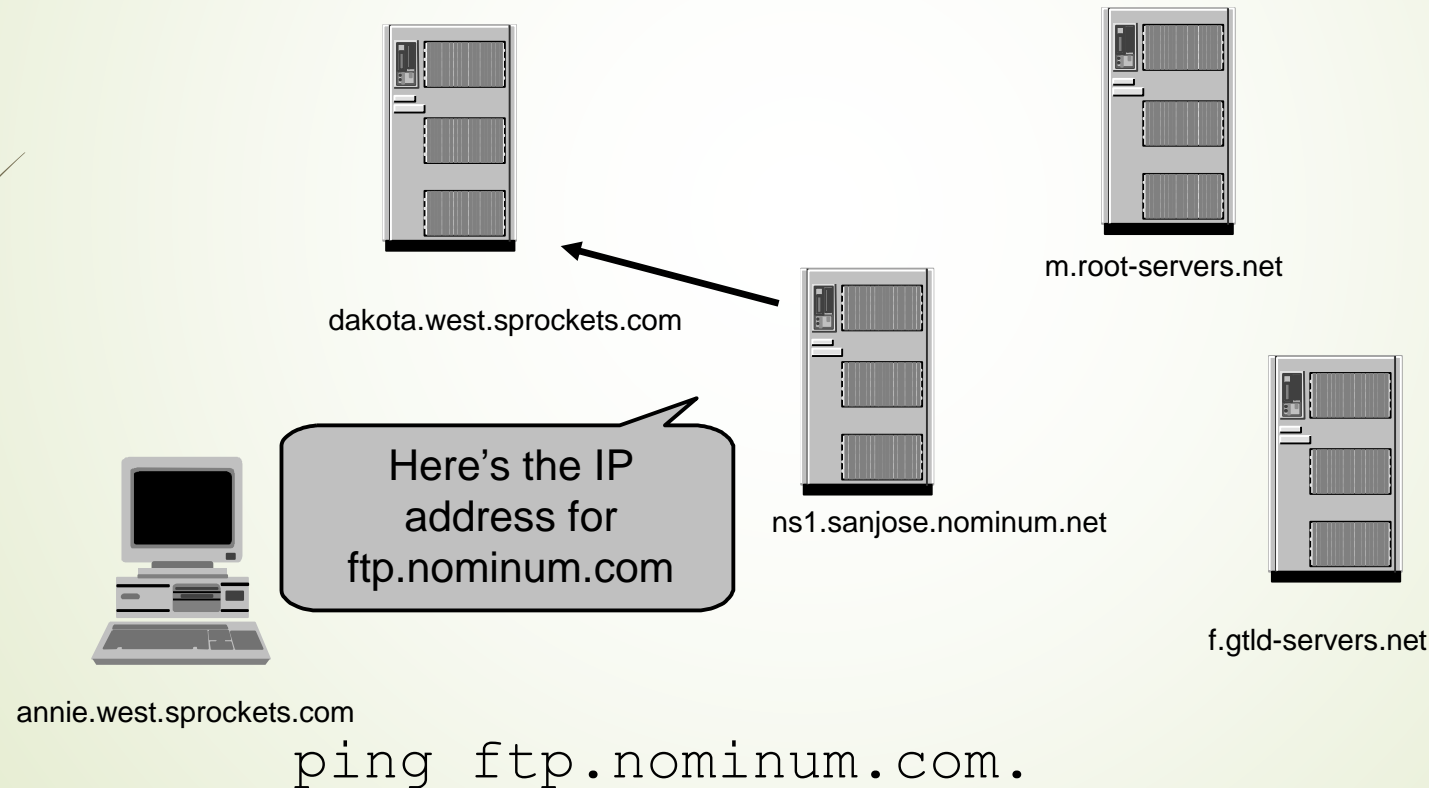
Resolution Process (Caching)

- ▶ *dakota* has **cached** a NS record indicating *ns1.sanjose* is an *nominum.com* name server, so it asks it for *ftp.nominum.com*'s address



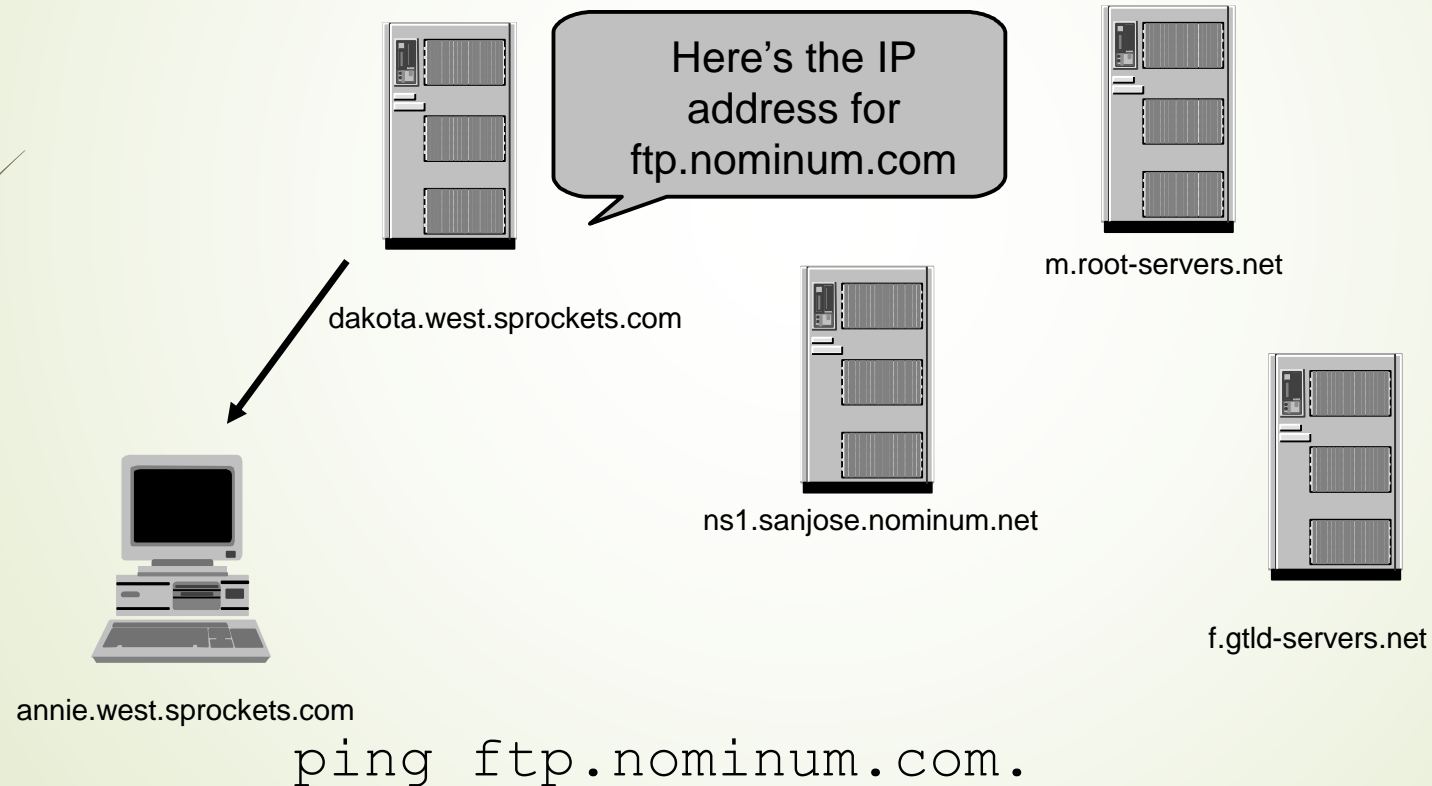
Resolution Process (Caching)

- The *nominum.com* name server *ns1.sanjose* responds with *ftp.nominum.com*'s address

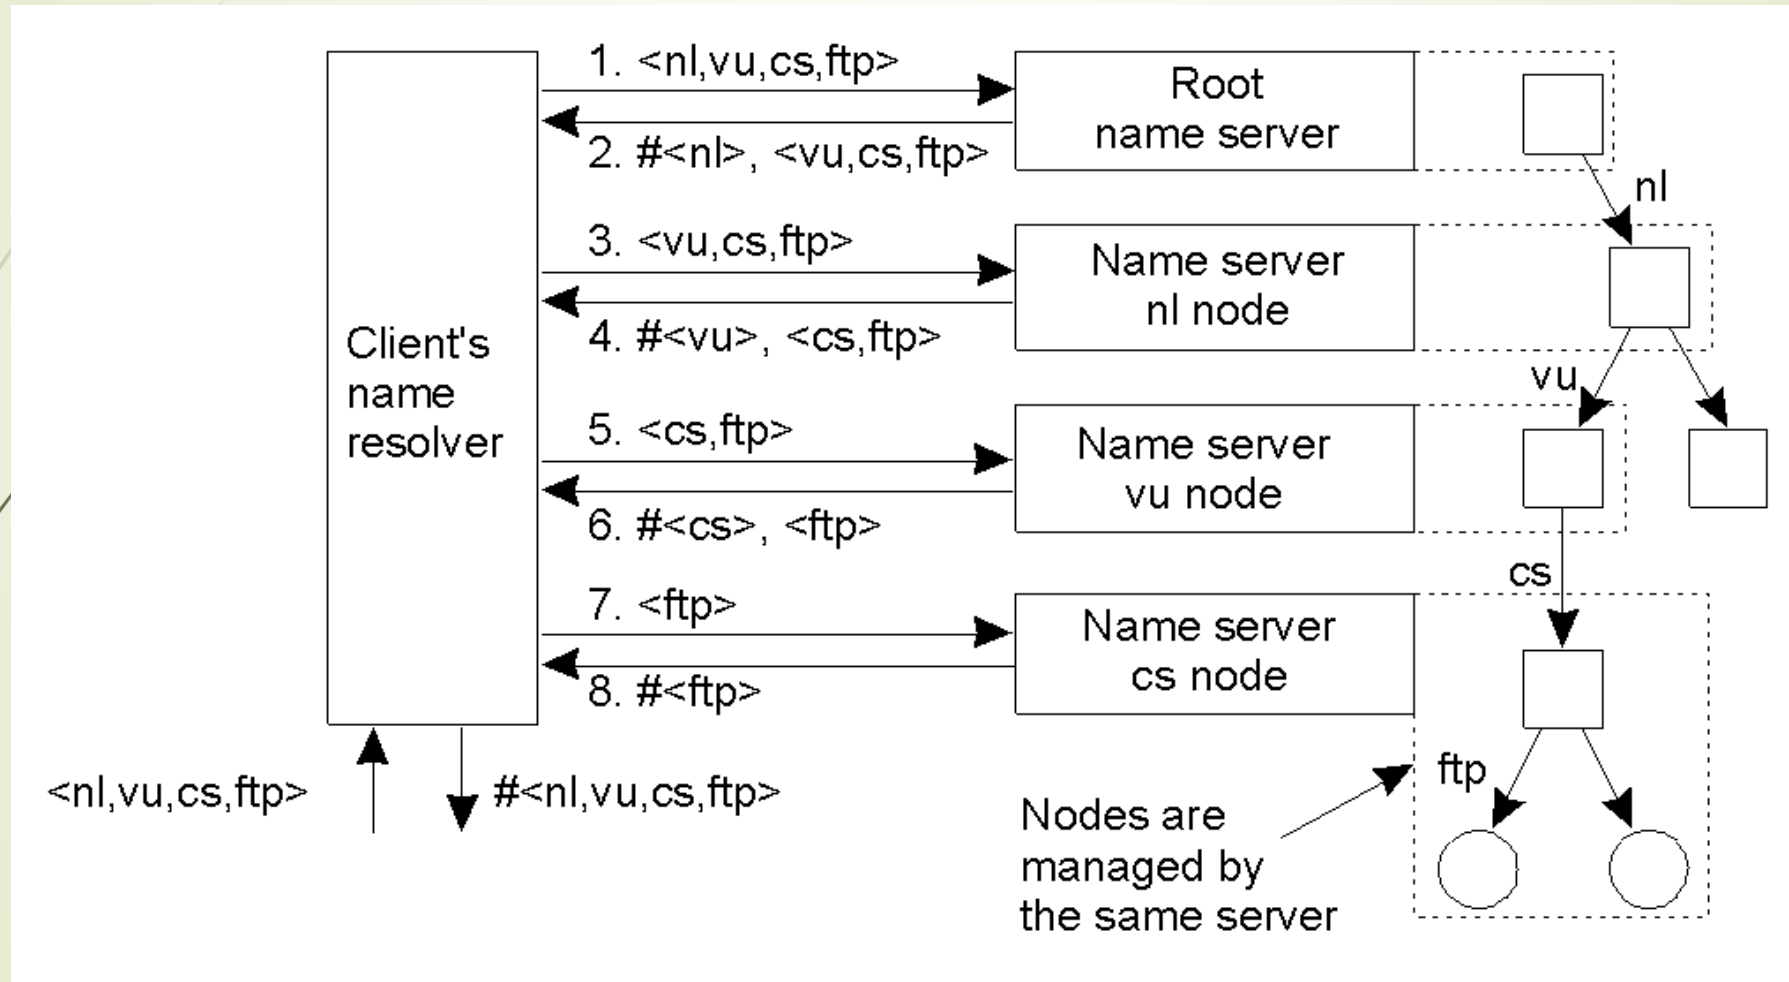


Resolution Process (Caching)

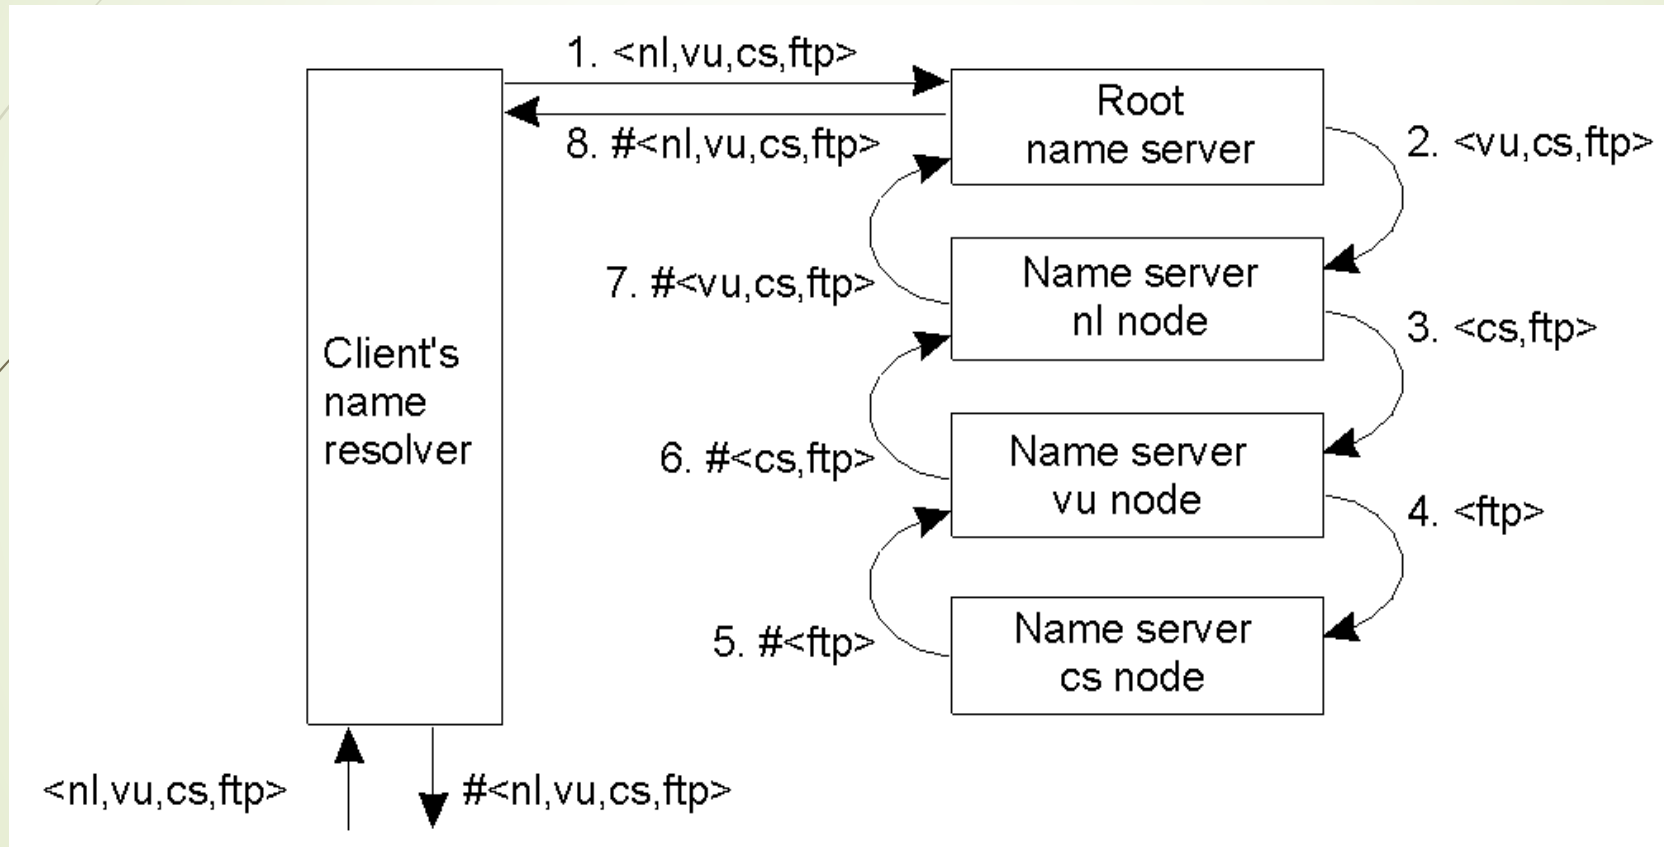
- ▶ The name server *dakota* responds to *annie* with *ftp.nominum.com*'s address



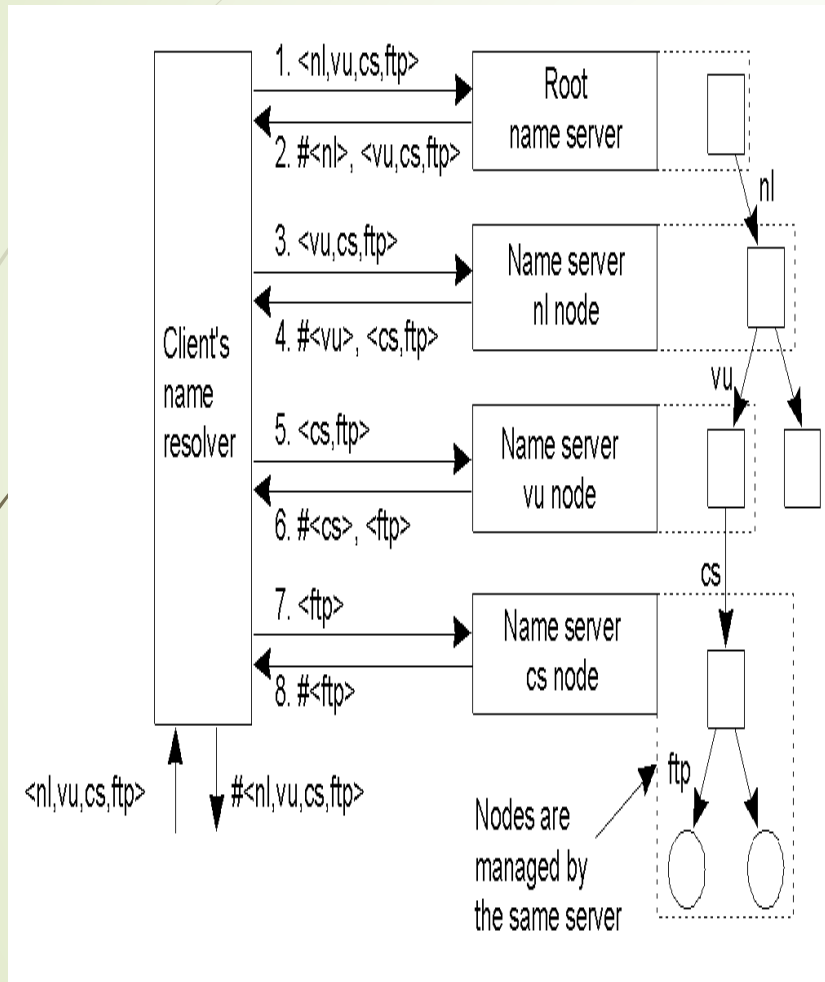
Iterative Name Resolution



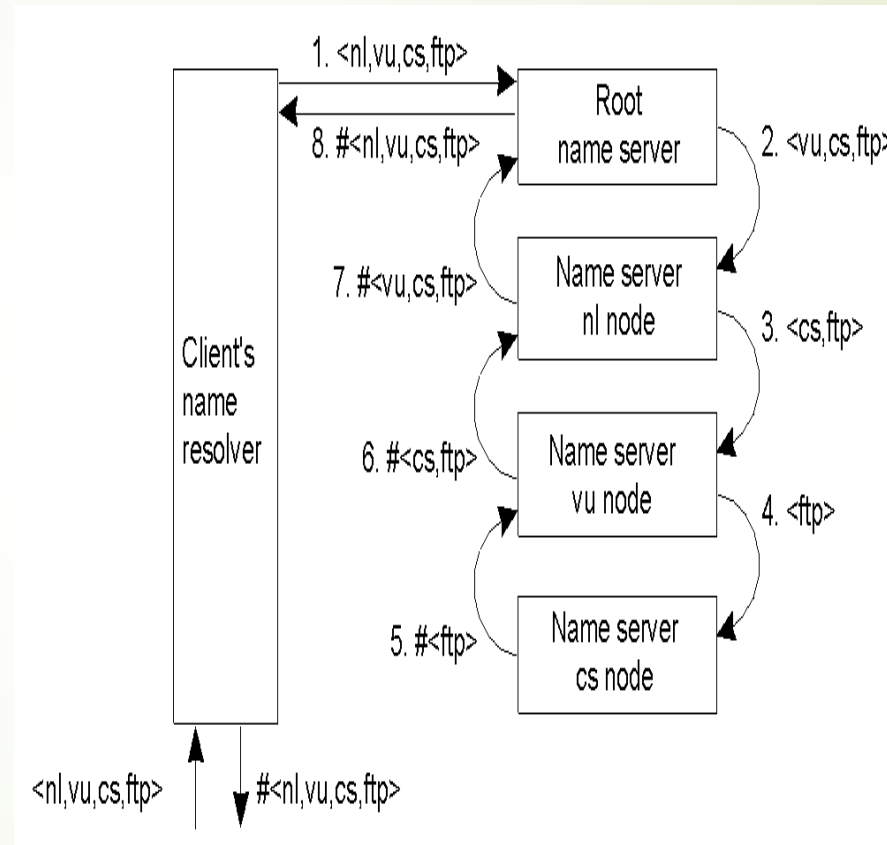
Recursive Name Resolution (1)



Iterative versus Recursive Resolution (1)




How about communication cost?

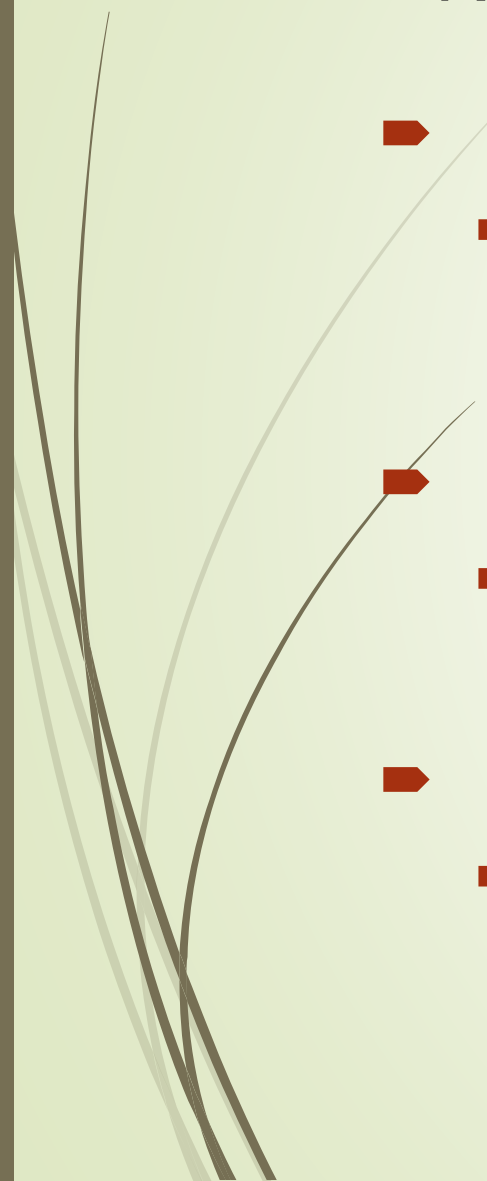


Performance-wise, which is better?

Which works better with caching?



Iterative versus Recursive Resolution (2)

- Performance-wise, which is better?
 - Recursive method puts higher performance demand on each name server
 - Which works better with caching?
 - Recursive method works better with caching
 - How about communication cost?
 - Recursive method can reduce communication cost
- 



Questions??

